Dear Hiring Manager,

I am applying for the AI Engineer role on DataDome's Threats Research team. The job is applied AI engineering in a specific sense: taking generative AI and agent orchestration out of a demo and into a detection pipeline that holds up under real traffic, which is what I have done for the past two and a half years.

At FlowSpace, I built and owned an AI pipeline that takes a workspace provider's brochure in any format, PDF, website, photo, or spreadsheet, and uses LLM APIs with agentic orchestration to extract structured data and draft a full listing, with a human reviewing before it reaches the backend. That pipeline has produced 900+ draft listings. I also built RAG-backed support agents that answer customer questions against live platform data, escalate to a human below a confidence threshold, and get evaluated against failure modes before reaching real users. Both run in production, not in a notebook.

DataDome's problem sits further along the same axis: telling legitimate traffic apart from bots and agents in milliseconds, across websites, APIs, and now MCP traffic, using generative AI to speed up how the Threats Research team investigates and responds. Using a crew of agents to automate traffic analysis for the SOC team is the same category of problem I have solved with LLM orchestration at FlowSpace, on a harder detection surface with higher stakes if it goes wrong.

I am based in Venice, Italy, and hold Italian EU long-term resident status, which authorises me to work in Italy but not automatically in France. For a France-based role I would need the employer-side permit process, and I would flag my preference to work remotely from Italy if that arrangement is possible. I am ready to relocate and start the permit process otherwise, and I am available immediately.

Regards,
Prabhavit Gupta